



BF16: TFLOPS vs matrix size (INNER_LOOP=10000)



BF16: efficiency vs matrix size



Same $Mt=Kt=T$, $Nt=8$ matmul shape, only fidelity differs. HiFi4 does 4x more real FPU work per instruction than LoFi, so it dilutes fixed perinstruction overhead as size grows and climbs toward its peak. LoFi's plateau stays flat because its bottleneck (likely) is FPU instruction issue-rate itself, not overhead that more size can dilute.